

# NetRush Protocol Mini-RFC

**Protocol:** NetRush v1.0 | **Signature:** NRSH | **Date:** December 24, 2025

---

## 1. Introduction

**NetRush** is a UDP-based real-time multiplayer game synchronization protocol for Grid Clash, a competitive grid capture game where 2-8 players claim cells on a  $5 \times 5$  grid.

### Why a new protocol?

- TCP: Head-of-line blocking unacceptable for real-time gameplay
- MQTT: Pub/sub model unsuitable for authoritative server architecture
- Existing protocols (ENet, RakNet): Prohibited by assignment requirements

### Design constraints:

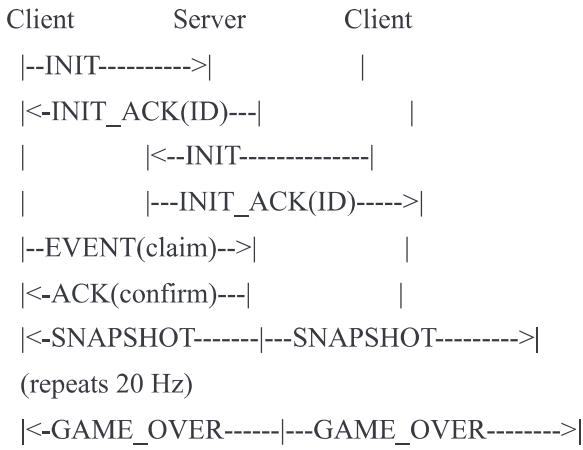
- Max packet: 1200 bytes (avoids fragmentation)
  - Update rate: 20 Hz (50ms intervals)
  - Acceptable loss: 0-30% for state updates
  - Transport: UDP over IPv4
  - No external game networking libraries
- 

## 2. Protocol Architecture

### Client-Server Model:

Entity	Role
Server	Authoritative state; processes claims; broadcasts 20 Hz snapshots
Client	Sends claim events; renders state with prediction/smoothing

### Communication Flow:



## Finite State Machines:

*Client:* DISCONNECTED → CONNECTING (send INIT, retry on timeout) → PLAYING (recv snapshots, send events) → GAME\_OVER

*Server per client:* IDLE → CONNECTED (recv INIT, assign ID, broadcast @ 20Hz) → GAME\_OVER (grid full, send winners 3×)

---

## 3. Message Formats

### 3.1 Binary Header (28 bytes)

Field	Offset	Size	Type	Description
protocol_id	0	4B	ASCII	"NRSH" signature
version	4	1B	uint8	Protocol version (1)
msg_type	5	1B	uint8	Message type (see §3.2)
snapshot_id	6	4B	uint32	Snapshot/sequence ID
seq_num	10	4B	uint32	Packet sequence number
timestamp_ms	14	8B	uint64	Unix epoch milliseconds
payload_len	22	2B	uint16	Payload bytes (max 1172)
checksum	24	4B	uint32	CRC32 of header+payload

**Struct format:** `!4sBBIIQHI` (network byte order)

## 3.2 Message Types

Value	Name	Direction	Purpose
0	INIT	C→S	Connect/keep-alive
1	INIT_ACK	S→C	Assign player ID
2	SNAPSHOT	S→C	State update (full/delta)
3	EVENT	C→S	Critical action (claim)
4	ACK	S→C	Confirm EVENT
5	GAME_OVER	S→C	Announce winners

## 3.3 Payload Formats (JSON, optional zlib compression)

**Compression flag:** First payload byte: `0x00`=uncompressed, `0x01`=zlib

**INIT:** `{}` (empty)

**INIT\_ACK:** `{"client_id": <int>}`

**SNAPSHOT:**

```
json
{
  "full": <bool>,
  "changes": [[row,col,owner], ...],
  "grid_enc": "r,c,owner;..." (if full),
  "redundant": [{"snapshot_id": <int>, "changes": [...]}]
}
```

**EVENT:** `{"cell": <int>, "client_id": <int>, "ts": <int>}`

**ACK:** `{"cell": <int>, "owner": <int>}`

**GAME\_OVER:** `{"winner": [<int>, ...], "final_grid_enc": "..."}`

**Grid encoding (sparse):** Only claimed cells: `"0,0,1;0,1,2"` = cell(0,0)→player1, cell(0,1)→player2

---

## 4. Communication Procedures

## Session Start:

1. Client sends INIT, waits 500ms for INIT\_ACK, retries up to 10×
2. Server assigns ID, sends INIT\_ACK, adds client to broadcast list
3. Client enters PLAYING state, listens for snapshots

## Normal Exchange:

1. Server broadcasts SNAPSHOT every 50ms (20 Hz):
  - Full grid every 10th snapshot (500ms)
  - Delta changes otherwise
  - Includes last 2 snapshots as redundancy
2. Clients receive, validate checksum, apply changes with deduplication (seen\_ids set)

## Critical Events (Claims):

1. Client clicks cell → sends EVENT, adds to retransmission queue (timeout: 500ms, max retries: 3)
2. Server processes claims (timestamp-ordered), sends ACK with confirmed owner
3. Client receives ACK → removes from queue, updates cell state

**Keep-Alive:** Client sends INIT every 3s (no explicit timeout enforcement)

**Termination:** Grid full → server computes winners → sends GAME\_OVER 3× (reliability) → stops broadcast

---

## 5. Reliability & Performance Features

### 5.1 Reliability Model

Data	Method
SNAPSHOT	Unreliable (next replaces lost)
EVENT	Reliable (timeout+retry, ACK)
INIT/INIT_ACK	Reliable (client retry)
GAME_OVER	Reliable (sent 3×)

## 5.2 Retransmission

**EVENT timeout:** Fixed RTO = 500ms (conservative for home networks, typical RTT 5-100ms)

**Max retries:** 3 attempts, then fail

## 5.3 Redundant Updates

**Mechanism:** Each SNAPSHOT includes changes from last K=2 snapshots

**Benefit:** Recovers from burst loss  $\leq 100\text{ms}$  (2 packets)

**Deduplication:** Clients track `(seen_ids)` (rolling window, max 500)

## 5.4 Delta Compression

- **Full snapshot:** Every 10th update (500ms), zlib compressed,  $\sim 100\text{-}300$  bytes
- **Delta snapshot:** Only changed cells, uncompressed,  $\sim 20\text{-}80$  bytes
- **Bandwidth savings:** 78% reduction vs. full snapshots every update

## 5.5 Sparse Grid Encoding

**Before:** 25 cells  $\times$  12 bytes = 300 bytes (dense JSON)

**After:** Only claimed cells: "r,c,owner;..." =  $\sim 60$  bytes (10 claims)

**Savings:** 80% for typical mid-game state

---

## 6. Experimental Evaluation Plan

### 6.1 Metrics

Metric	Method
Latency	$\text{server\_ts} - \text{client\_recv\_ts}$
Jitter	$\sigma(\text{inter-packet arrival})$
Loss rate	gaps in <code>snapshot_id</code> / expected
Bandwidth	bytes/sec, logged by server
CPU	<code>psutil.cpu_percent()</code>

### 6.2 Test Scenarios

**Baseline:** Localhost, 0% loss,  $< 5\text{ms}$  RTT

**Moderate Loss:** 10% random loss via `tc qdisc add dev lo root netem loss 10%`

**High Latency:** 150ms delay via `tc qdisc add dev lo root netem delay 150ms 20ms`

**Burst Loss:** 20% loss w/ 50% correlation via `tc qdisc add dev lo root netem loss 20% 50%`

**Combined:** 100ms delay + 15% loss + 2Mbps limit:

```
bash
```

```
tc qdisc add dev lo root handle 1: tbf rate 2mbit burst 32k latency 400ms
```

```
tc qdisc add dev lo parent 1:1 netem delay 100ms 10ms loss 15%
```

## 6.3 Automated Testing

**Script:** `run_experiment.sh <scenario> <duration>`

```
bash
```

```
#!/bin/bash
```

```
SCENARIO=$1; DURATION=${2:-60}
```

```
# Setup netem (varies by scenario)
```

```
python3 server.py &
```

```
sleep 2
```

```
python3 client.py & python3 client.py &
```

```
sleep $DURATION
```

```
killall python3
```

```
tc qdisc del dev lo root
```

```
mv *_log.csv results/$SCENARIO/
```

**Analysis:** `analyze_results.py` loads CSVs, computes metrics, generates plots (latency over time, bandwidth)

## 6.4 Expected Results

Scenario	Latency	Loss	Outcome
Baseline	1-5ms	0%	Smooth gameplay
Moderate Loss	5-10ms	~10%	Redundancy compensates
High Latency	145-175ms	0%	Playable, claims confirm <700ms
Burst Loss	5-10ms	~20%	Smoothing masks loss
Combined	90-110ms	~15%	Functional, occasional artifacts

## 7. Example Use Case Walkthrough

**Scenario:** Player 1 claims cell(0,0), Player 2 claims cell(0,1)

Time	Event	Packet
0ms	P1 starts	-
5ms	P1→INIT	INIT(sid=0,seq=0)
8ms	S→P1	INIT_ACK(cid=1)
175ms	P1 clicks (0,0)	-
176ms	P1→EVENT	EVENT(cell=0,cid=1,ts=176)
200ms	Server processes	Grid[0][0]=1
200ms	S→P1	ACK(cell=0,owner=1,seq=3)
200ms	S→All	SNAPSHOT(sid=3,changes=[[0,0,1]])

### Packet hexdump (INIT):

```
4e 52 53 48 01 00 00 00 00 00 00 00 00 00 00 00  
00 00 00 00 05 00 02 00 7b 7d xx xx xx xx
```

- `4e52 5348` = "NRSH"
- `01` = version, `00` = INIT type
- `7b7d` = JSON `{}`
- `xxxxxxxx` = CRC32 checksum

## 8. Limitations & Future Work

### Current Limitations:

1. No authentication (vulnerable to spoofing) → add shared secret in INIT
2. Fixed RTO=500ms (suboptimal for varied networks) → implement adaptive RTO (Karn's algorithm)

3. No encryption (plaintext) → add DTLS or AES
4. Single-threaded server → multi-threaded claim processing
5. Hardcoded  $5 \times 5$  grid → negotiate size in INIT\_ACK
6. Redundancy K=2 only → adaptive based on measured loss
7. No reconnection logic → add session tokens

### Future Enhancements:

- Adaptive redundancy:  $K = f(\text{loss\_rate})$
  - Client prediction: extrapolate positions during loss
  - Forward error correction (FEC): Reed-Solomon codes
  - Congestion control: monitor RTT, adjust update rate
  - NAT traversal: STUN/TURN integration
  - Anti-cheat: server-side timestamp validation, rate limiting
- 

## 9. References

1. RFC 768 - User Datagram Protocol (<https://www.rfc-editor.org/rfc/rfc768>)
2. RFC 6298 - Computing TCP's Retransmission Timer (timeout reference)
3. Gaffer on Games - "UDP vs TCP" ([https://gafferongames.com/post/udp\\_vs\\_tcp/](https://gafferongames.com/post/udp_vs_tcp/))
4. Valve Source Multiplayer Networking (delta compression techniques)
5. Python struct module documentation (<https://docs.python.org/3/library/struct.html>)
6. Python zlib module documentation (<https://docs.python.org/3/library/zlib.html>)